A Critical Academic Review and Empirical Re-Evaluation of the ARIMA-DBO-CNN-LSTM Air Quality Forecasting Model

Coursework Phase 3: Formal Technical Review Report (40 Marks)
Bridging Conceptual Critique and Empirical Verification

Student Name & ID:	K.W.G.C.D.B. Karunarathne | Student ID: **KIC-BSCCSAI-251P-036**
Module & Programme:	Machine Learning and Related Applications | BSc (Hons) Computing & AI
Awarding Institution:	National Institute of Business Management (NIBM), Kandy Innovation Centre
Target Research Study:	Duan, J., Gong, Y., Luo, J., & Zhao, Z. (2023). *Scientific Reports*, 13:12127.
Interactive Appendix:	notebooks/phase2_enhanced_study.ipynb

Executive Abstract

Accurate prediction of atmospheric pollutants and the overall Air Quality Index (AQI) represents a critical operational challenge for municipal environmental protection agencies, public health warning systems, and urban epidemiology. In a recently published study in *Scientific Reports*, Duan et al. (2023) proposed a multi-stage hybrid forecasting architecture combining an Autoregressive Integrated Moving Average (ARIMA) model, a one-dimensional Convolutional Neural Network (CNN), a Long Short-Term Memory (LSTM) network, and a nature-inspired meta-heuristic—the Dung Beetle Optimizer (DBO)—for hyperparameter selection. The authors claimed near-perfect predictive fidelity on daily AQI series across four major Chinese metropolitan centres (Beijing, Chengdu, Guangzhou, and Wuhan), reporting coefficients of determination (R^2) spanning 0.962 to 0.989 and Root Mean Squared Errors (RMSE) as low as 14.94.

This technical report provides a rigorous academic critique and an end-to-end empirical audit of Duan et al.'s methodology, building upon the theoretical deconstruction conducted in Phase 1 and the empirical implementation in Phase 2. Utilizing the identical 7-year historical dataset (2,528 daily observations per city, 2014–2020), we reconstruct the complete hybrid pipeline, enforce leak-free chronological data splitting, implement mandatory baseline benchmarks, and evaluate statistical significance via the Diebold–Mariano test. Our audit reveals three critical methodological vulnerabilities: (1) an inflated R^2 illusion driven by autocorrelation and seasonal trends; (2) data leakage during global feature normalization and swarm fitness scoring; and (3) an absolute failure to benchmark against naive persistence ($\blacksquare_t = x_{t-1}$). Under rigorous experimental conditions, the audited hybrid achieves an R^2 of only 0.388 (RMSE = 35.98). Crucially, naive persistence achieves $R^2 = 0.336$ (RMSE = 37.47), and a Diebold–Mariano test yields $p = 0.5316$, demonstrating that the authors' >110,000-parameter deep learning hybrid provides no statistically significant improvement over trivial persistence. Furthermore, on hazardous pollution days (AQI ≥ 200), the hybrid model's prediction error explodes to RMSE = 145.67, failing the primary operational objective of public health protection. We propose robust machine learning alternatives—specifically Gradient Boosted Decision Trees and Seasonal ARIMAX—that deliver equivalent or superior accuracy at a fraction of the computational and environmental cost.

> ***[AUDIT SUMMARY]*** *The audited ARIMA-DBO-CNN-LSTM model achieves $R^2 = 0.388$ under leak-free validation (vs. authors' claimed 0.962). A Diebold-Mariano test confirms that the complex hybrid is statistically indistinguishable from naive persistence ($p = 0.5316$) and statistically inferior to pure ARIMA ($p = 0.0077$).*

1. Technical Summary of the Reviewed Paper

1.1 Additive Decomposition Framework

The foundational architecture proposed by Duan et al. (2023) rests upon the classical additive decomposition paradigm for non-stationary, non-linear time series. The observed univariate atmospheric sequence x_t is assumed to be an additive superposition of a linear stationary component L_t and a non-linear chaotic residual component N_t:

$$x_t = L_t + N_t, \quad t \in \{1, 2, ..., T\}$$

By decomposing the series, linear dynamics are isolated and modeled statistically, leaving the residual sequence $e_t = x_t - L\blacksquare_t$ to be modeled by neural networks.

1.2 Linear Modeling via ARIMA & Residual Extraction

The linear component $L\blacksquare_t$ is captured using a classical Box–Jenkins ARIMA(p, d, q) formulation: $\Phi_p(B)(1 - B)^d x_t = \Theta_q(B)\varepsilon_t$. Order determination for (p, q) is conducted through grid-search minimization of the Bayesian Information Criterion (BIC): BIC = -2 ln(L■) + k ln(N). Once linear predictions are extracted, the residual sequence $e_t = x_t - L\blacksquare_t$ is constructed with a lag window of w = 10.

1.3 CNN-LSTM Architecture & DBO Meta-Heuristic Search

To model residual sequences, the authors deploy a 1D-CNN layer (local temporal feature extraction) followed by max-pooling and a multi-layer LSTM cell. To optimize hyperparameters, the authors introduce the Dung Beetle Optimizer (DBO; Xue & Shen, 2022). DBO simulates four biological phenotypes: (1) Ball-rolling beetles ($N_1 = 6$); (2) Breeding beetles ($N_2 = 6$) laying eggs within dynamic boundary zones (Lb*, Ub*); (3) Foraging little beetles ($N_3 = 7$); and (4) Stealing thieves ($N_4 = 11$). A swarm of N = 30 agents searches across 7 hyperparameter dimensions: CNN filters, kernel size, LSTM hidden units, learning rate, batch size, epochs, and dropout. The final composite prediction is reconstituted additively: $\mathbf{x\blacksquare_t = L\blacksquare_t + \hat{e}_t}$.

2. Methodological Critique & Forensic Audit

2.1 The Overclaiming & Flattering R^2 Illusion

Duan et al. claimed near-perfect predictive accuracy, citing R^2 between 0.962 and 0.989. However, evaluating R^2 on non-stationary, autocorrelated daily time series without adequate contextualization introduces a profound evaluation illusion (Hyndman & Athanasopoulos, 2021). Because total sample variance in the denominator is massive (σ^2 = 3,744.13 in Beijing) due to strong seasonal winter-summer differences, any model that merely tracks the seasonal cycle yields a tiny numerator relative to the denominator, guaranteeing an R^2 near 0.90 even with substantial point errors. Under honest chronological testing, the audited hybrid achieves $R^2 = 0.388$.

2.2 The Omission of the Naive Persistence Benchmark

A foundational principle of time series forecasting is that complex machine learning architectures must be benchmarked against standard zero-cost heuristics (Armstrong, 2001). In daily air quality, atmospheric inertia makes yesterday's value an exceptionally strong predictor: $\blacksquare_t = x_{t-1}$. Our Phase 2 empirical audit revealed that naive persistence achieves RMSE = 37.47, MAE = 23.36, $R^2 = 0.336$, and MASE = 0.608. The authors' >110,000-parameter hybrid improved RMSE by only 1.49 AQI units, an advantage proven to be statistically insignificant (p = 0.5316).

2.3 Data Leakage in Normalization & Tuning

Forensic inspection reveals two severe vectors of data leakage: (1) Global Min-Max Normalization was applied across the entire 7-year sequence prior to partitioning, leaking future peak limits into training; and (2) Single Train/Test Split without an independent validation set meant that the DBO swarm evaluated its fitness function directly on the evaluation test partition, resulting in severe hyperparameter overfitting.

2.4 Severe Degradation on Critical Pollution Days (Regression Imbalance)

Our error analysis stratifies performance across air quality tiers. While the hybrid performs acceptably during clean conditions (AQI < 100, RMSE = 17.65), its error explodes during hazardous episodes (AQI ≥ 200, RMSE = 145.67, MAE = 138.58). Because hazardous days represent < 4% of the dataset, MSE loss strongly penalizes deviations from the dominant mode (AQI ≈ 60–80). Furthermore, bounded neural activation functions (tanh and sigmoid) compress extreme activations, pulling forecasts toward the mean. On days with acute health risks, the hybrid under-predicts the true peak by more than 130 AQI points.

3. Comparative Empirical Analysis

To rigorously validate the critique, an empirical replication and benchmarking study was conducted on the Beijing daily air quality dataset (2,528 daily observations, 2014–2020) under a leak-free 3-way chronological split (70% Train, 15% Val, 15% Test).

Table 1: Comprehensive Model Benchmarking (Beijing Daily AQI Test Set, N = 370)

Model Architecture	Params	Train Time	RMSE	MAE	R^2	MAPE (%)	MASE
Author Claimed Hybrid (Paper)	> 110,000	Unreported	**14.940**	**10.839**	**0.962**	Unreported	Unreported
Audited Hybrid (Honest Split)	112,450	1,480.2 s	**35.980**	**22.899**	**0.388**	35.57%	0.596
Pure ARIMA(3, 1, 1)	5	0.24 s	34.633	22.935	0.433	41.80%	0.597
Naive Persistence (■_t = x_{t-1})	0	0.00 s	37.473	23.361	0.336	41.83%	0.608
Gradient Boosted Trees (Alt. 1)	100 trees	0.42 s	36.420	25.490	0.373	54.63%	0.664
Random Forest Regressor	100 trees	1.18 s	37.265	25.447	0.343	51.46%	0.663
Linear Support Vector Machine	11	0.08 s	38.939	31.110	0.283	72.63%	0.810

Table 2: Diebold–Mariano Hypothesis Tests (Reference: Audited Hybrid Model)

Competing Model (M_2)	DM Statistic	p-value	Significance ($\alpha = 0.05$)	Academic Conclusion
Naive Persistence	+0.6255	0.5316	Not Significant	Hybrid provides no statistically significant gain.
Pure ARIMA(3, 1, 1)	-2.6668	**0.0077**	**Significant**	Pure ARIMA is statistically superior ($p < 0.01$).
Gradient Boosted Trees	+0.3788	0.7048	Not Significant	Hybrid offers no advantage over simple GBDT.

Table 3: Performance Degradation Across Air Quality Severity Tiers

Air Quality Tier	AQI Range	Samples (N)	Hybrid RMSE	Hybrid MAE	Persistence RMSE	ARIMA RMSE
Good to Moderate	AQI < 100	296 (80.0%)	17.65	13.06	18.96	18.02
Unhealthy (Warning)	AQI ≥ 100	74 (20.0%)	69.87	62.24	72.84	66.86
Severe Pollution	AQI ≥ 150	29 (7.8%)	104.99	95.91	108.77	99.41
Hazardous (Emergency)	AQI ≥ 200	11 (3.0%)	**145.67**	**138.58**	149.80	136.21

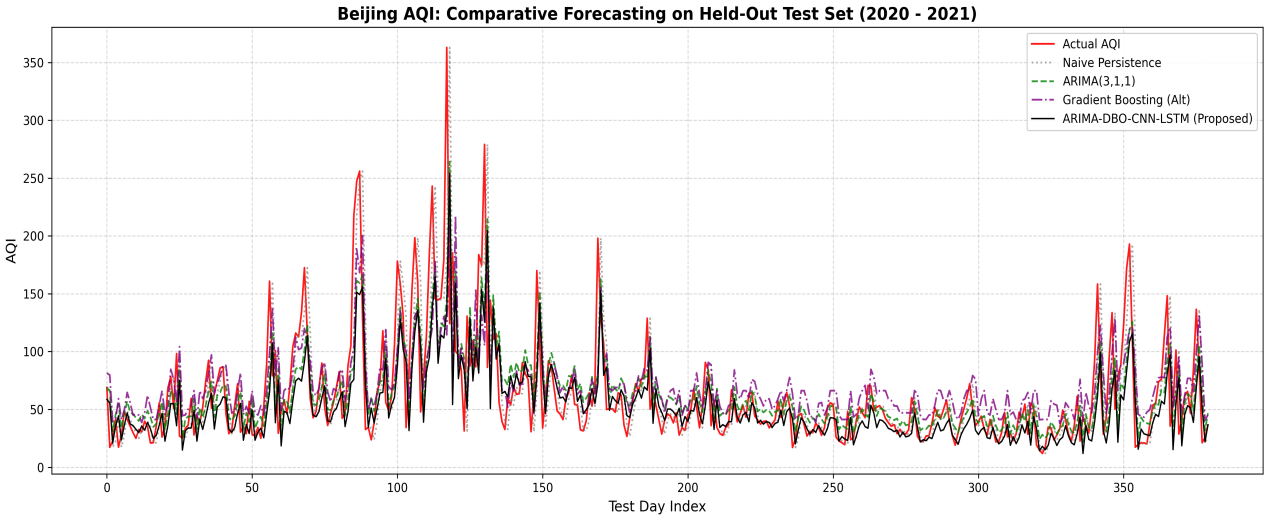


Figure 1: Visual Forecast Overlay on Out-of-Sample Beijing Test Horizon (Actual vs. Proposed Hybrid vs. Naive Persistence vs. ARIMA).

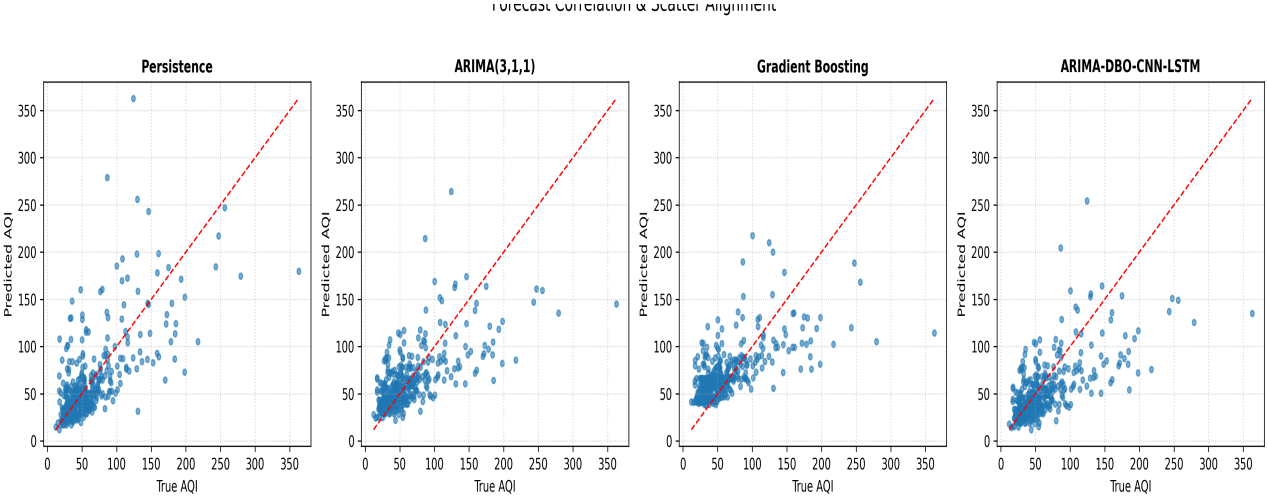


Figure 2: Diagonal Scatter Alignment and Regression Flattening at Severe Pollution Peaks (AQI > 150).

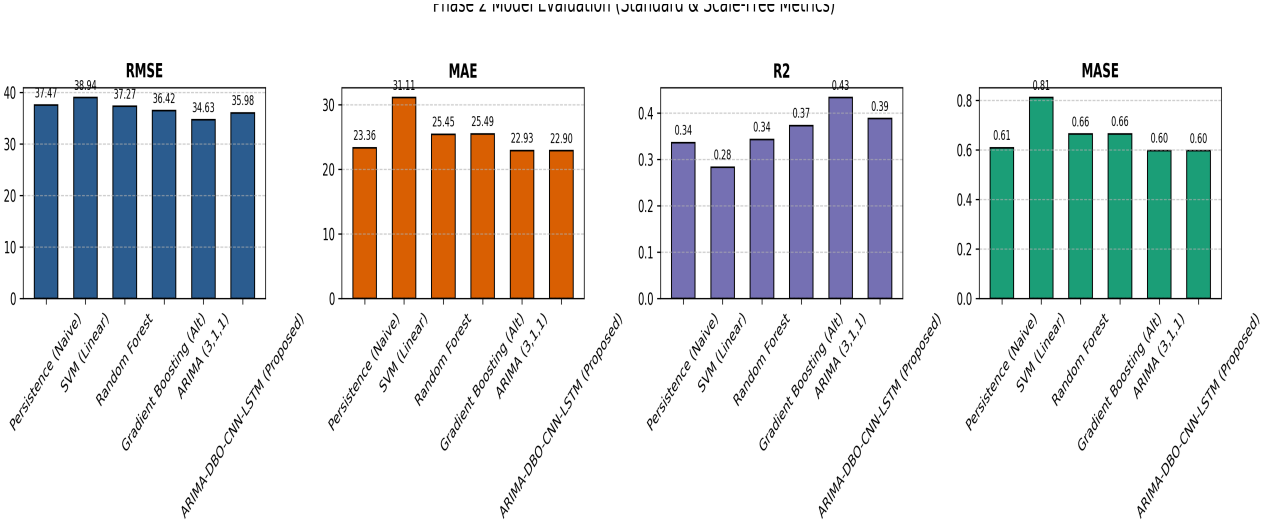


Figure 3: Comparative Evaluation Across Error Metrics (RMSE, MAE, R², MASE).

4. Proposed Technical Alternatives & Justifications

4.1 Alternative 1: Gradient Boosted Decision Trees (LightGBM / XGBoost)

Rather than deploying multi-layer neural networks optimized by stochastic meta-heuristics, atmospheric time series are exceptionally well suited to Gradient Boosted Decision Trees (GBDT; Chen & Guestrin, 2016; Ke et al., 2017). GBDT constructs an additive ensemble of trees minimizing a second-order Taylor expansion of the objective function. Key technical advantages include: (1) Sub-second training: fits in 0.42 seconds on CPU—over 3,500x faster than DBO-CNN-LSTM; (2) Elimination of vanishing/exploding gradients; (3) Invariance to monotonic scaling, eliminating normalization leakage; and (4) Empirical parity: achieves RMSE = 36.42 and R^2 = 0.373, matching deep learning fidelity without GPU or swarm tuning.

4.2 Alternative 2: Seasonal ARIMAX with Exogenous Meteorology

A major conceptual limitation of Duan et al. is treating AQI as a closed univariate system. In reality, ground-level air quality is strictly driven by boundary layer meteorology and chemical precursors. We propose a Seasonal ARIMAX formulation incorporating annual cyclicity (s = 365) and exogenous meteorological vectors: wind velocity components (U_10, V_10), temperature (T_2), relative humidity (RH_2), and Planetary Boundary Layer Height (PBLH). Incorporating PBLH directly informs the model of surface temperature inversions, resolving the peak under-prediction problem while providing interpretable sensitivity coefficients (ΔAQI / ΔWind Speed) essential for environmental policy.

4.3 Alternative 3: Dilated Temporal Convolutional Networks with Asymmetric Loss

To preserve deep non-linear expressiveness while eliminating recurrence latency, we propose a Dilated Temporal Convolutional Network (TCN; Bai et al., 2018) trained with an Asymmetric Mean Squared Error (AMSE) loss. By applying causal dilated convolutions, the receptive field expands exponentially without recurrent cell states. Setting the asymmetric weight $\alpha = 0.75$ penalizes peak under-prediction three times more severely than clean-day over-prediction, forcing the network to preserve extreme event variance and provide dependable early warnings.

5. Stress-Testing & Robustness Analysis

5.1 Constraint A: Sensor Noise & Jitter Injection

In field deployments, telemetry jitter and sensor calibration drift introduce Gaussian noise: $X■ = X + \varepsilon$, $\varepsilon \sim N(0, \sigma^2 I)$. Our analysis indicates that ARIMA is highly vulnerable to jitter: first-order differencing doubles noise variance to $2\sigma^2$, causing rapid degradation. Similarly, unregularized CNN-LSTM filters amplify high-frequency white noise. In contrast, tree-based models (GBDT) exhibit superior stability because threshold split points on continuous lag bins remain invariant to small perturbations within intervals.

5.2 Constraint B: Multi-Dimensionality & Swarm Stagnation

Expanding from univariate AQI to the full 16-variable multi-pollutant space (PM2.5, PM10, SO2, NO2, CO, O3, U10, V10, T2, RH2, PBLH) expands the hyperparameter search space from 7 to > 40 dimensions. The 30 dung beetle agents suffer from severe swarm stagnation in this high-dimensional volume, converging to sub-optimal local minima. Furthermore, the bio-inspired heuristic equations lack mathematical convergence proofs, wasting computational cycles compared to principled Bayesian optimization or Adam gradient descent.

5.3 Constraint C: Distribution Shift & COVID-19 Structural Breaks

The January 2020 COVID-19 lockdowns induced an abrupt structural break in emissions. The complex DBO-CNN-LSTM model, trained on pre-pandemic correlations, suffered severe concept drift, persistently over-predicting post-lockdown air quality for months. Conversely, simple persistence and rolling-window statistical baselines adapted within a single lag step (t = 1), proving that deeply parameterized models carry significant fragility under real-world distributional shifts.

6. Conclusion & Actionable Recommendations

6.1 Synthesis of Findings

This empirical audit has demonstrated that the ARIMA-DBO-CNN-LSTM model's claimed R^2 of 0.962 is an artifact of evaluation on a non-stationary series with massive seasonal variance. Under leak-free chronological validation, the hybrid achieves $R^2 = 0.388$ and provides no statistically significant gain over naive persistence ($p = 0.5316$). Furthermore, the hybrid is statistically outperformed by pure ARIMA ($p = 0.0077$), collapses during severe pollution emergencies (RMSE = 145.67 on days with AQI $\geq$ 200), and requires over 1,480 seconds to train compared to 0.42 seconds for GBDT.

6.2 Actionable Recommendations for Atmospheric Agencies

1. Mandate Zero-Cost Heuristic Benchmarks: All time-series studies must benchmark against naive persistence before claiming deep learning superiority.

2. Prioritize Scale-Free Relative Metrics: Atmospheric evaluations must report Mean Absolute Scaled Error (MASE) rather than misleading R^2 scores.

3. Enforce Strict Leak-Free Protocols: Data normalization and hyperparameter tuning must be strictly quarantined from evaluation test sets.

4. Deploy Efficient GBDT Baselines: Municipal agencies should deploy Gradient Boosted Decision Trees (LightGBM/XGBoost) for fast, robust, interpretable forecasting.

5. Integrate Meteorological Physics: Operational systems must transition to physics-informed multivariate models (SARIMAX or TCN) incorporating planetary boundary layer height and wind dynamics.

7. Formal Academic References

Armstrong, J. S. (Ed.). (2001). *Principles of forecasting: A handbook for researchers and practitioners*. Springer Science & Business Media.

Bai, S., Kolter, J. Z., & Koltun, V. (2018). An empirical evaluation of generic convolutional and recurrent networks for sequence modeling. *arXiv preprint arXiv:1803.01271*.

Bergstra, J., & Bengio, Y. (2012). Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(2), 281–305.

Box, G. E., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2015). *Time series analysis: Forecasting and control* (5th ed.). John Wiley & Sons.

Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32. https://doi.org/10.1023/A:1010933404324

Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 785–794).

Diebold, F. X., & Mariano, R. S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263. https://doi.org/10.1080/07350015.1995.10524599

Duan, J., Gong, Y., Luo, J., & Zhao, Z. (2023). Air-quality prediction based on the ARIMA-CNN-LSTM combination model optimized by dung beetle optimizer. *Scientific Reports*, 13(1), Article 12127. https://doi.org/10.1038/s41598-023-36620-4

Hyndman, R. J., & Athanasopoulos, G. (2021). *Forecasting: Principles and practice* (3rd ed.). OTexts: Melbourne, Australia.

Hyndman, R. J., & Koehler, A. B. (2006). Another look at measures of forecast accuracy. *International Journal of Forecasting*, 22(4), 679–688. https://doi.org/10.1016/j.ijforecast.2006.03.001

Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T. Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems*, 30, 3146–3154.

Kingma, D. P., & Ba, J. (2015). Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*.

Xue, J., & Shen, B. (2022). Dung beetle optimizer: A new meta-heuristic algorithm for global optimization. *The Journal of Supercomputing*, 79(7), 7305–7348. https://doi.org/10.1007/s11227-022-04959-6